

**DUE: Wednesday, February 13, 2:20PM in class.**

You may discuss this assignment with whomever you wish, but please submit work in teams of *two to three* people. In particular, teams are encouraged to discuss the assignment with other teams, but please submit one assignment per team. Please, no individual submissions. Please indicate all team member names on the submission. Please prepare solutions in a *concise and organized fashion*. I prefer typeset presentations (e.g., cut and paste code/output into MS Word with added exposition when appropriate; knitr/R Markdown via Rstudio, etc.)—probably most appropriate for presenting (fixed width font) code/output, at least—but neatly handwritten presentations may also be appropriate for some problems. Sloppily prepared/disorganized solutions will not receive full credit.

**NOTE:** You will be asked to provide peer assessment for each of your team members (excluding yourself!) via a score from 0 to 10, 10 being best. Share these assessments only with me, please. Please use a standard 8.5x11 inch sheet of paper for your scores.

For this homework, we will work with data that I obtained from [MR07, Chap. 3] (who obtained it elsewhere). The data consist of measurements associated with processionary caterpillars. The observed output,  $y$ , is the (natural log of the) average number of caterpillar nests per tree in 500  $m^2$  of forest, and is in the *last* column (11) of the data set. The first ten columns consist of corresponding observations of  $k = 10$  input variables, roughly described in the following table.

$x_1$	elevation ( $m$ )
$x_2$	slope (degrees)
$x_3$	number of pines in the area
$x_4$	representative tree height ( $m$ )
$x_5$	representative tree diameter
$x_6$	settlement density index
$x_7$	site orientation index
$x_8$	height of dominant tree ( $m$ )
$x_9$	vegetation strata index
$x_{10}$	settlement mix index

NOTE: You may have to download/install an R package, which I assume you can do.

1. The data are available in BbLearn in the file, `caterpillar.txt`. Read the caterpillar data set into R and use R to: (i) name the columns of the resulting data frame as  $y$ ,  $x_1, \dots, x_{10}$ , and (ii) display the first 3 and last 3 records of your data set. Show your code and output.
2. Create an appropriate scatter plot matrix, similar to that seen in our notes. You may use the `pairs` function or a similar function in R. Use the `labels` option of the `pairs` function to provide more informative variable names than given above. If you use another plot function, label your plot in a similar manner. Show your code and plot.

3. How many linear models are possible using all possible subsets of the inputs, including the model consisting only of an intercept?
4. Use the function, `regsubsets`, in the `leaps` package, to perform an exhaustive search for the best model/subset. In particular, create a matrix of plots, one for each of  $SSE$ ,  $R^2$ ,  $R_a^2$ ,  $C_P$ , and  $BIC$ , each plot showing the *single* best model for each of the possible sizes. For the plot of  $C_P$  vs. number of covariates, add the reference line,  $C_P = 1 + |P|$ . Show your code and plots, appropriately annotated.
5. Identify the best model(s), over all sizes, as selected by each of the criteria,  $BIC$ ,  $C_P$ , and  $R_a^2$ ; this will result in 1, 2, or 3 models, depending on how these criteria (dis)agree.
6. Why does it not seem plausible to use the validation set approach, using MSPR, for this data set? Please be concise.
7. We only briefly introduced  $K$ -fold cross-validation in our notes as a more popular and modern (ML) alternative to the validation set approach to estimate generalization error and select models. Here, we conduct  $K$ -fold CV using  $K = n$  ( $n$ -fold CV, or CV( $n$ ), or LOOCV). For this homework, we perform  $n$ -fold CV. That is,  $K = n$  folds produce  $n$  training sets of size  $n_i = n - 1$  to predict on each of the corresponding  $n$  validation sets of size  $n_i^* = 1$ .

We'll leave the details to the function, `cv.glm`, in the `boot` package. To use this function, we need to fit our model(s) using the function, `glm`, in the `stats` package. For example, the code, below, uses `glm` to fit a linear model, just as `lm` does, and just as `regsubsets` does, above. The code then performs  $n$ -fold CV ( $K = n$  by default) using `cv.glm`. CV( $K=n$ ), as we have defined it, is contained in the first element of the list component named `delta` in the list object, `eg.cv.n`, returned by `cv.glm`.

```
> eg.glm<- glm(y ~ x1 + x2 + x3, data=cat.df)
> library(boot)
> eg.cv.n<- cv.glm(data=cat.df, glmfit=eg.glm)
> (cv.n<- eg.cv.n$delta[1]) ## CV(n)

[1] 0.426571
```

Using the above code to get you started, perform CV( $n$ ) on the best fitted model of each size as indicated by the `regsubsets` procedure, above. You will have to figure out what these models are (you should already know 1-3 of these, by previous work), fit each one using `glm`, then compute CV( $n$ ) for each one. Prepare a plot of CV( $n$ ) vs model size, similar to the above plots using other criteria. (NOTE: We just consider these few fitted models for which to compute CV( $n$ ); we do not compute CV( $n$ ) for all the models that you counted in problem 3.)

8. Discuss your findings. In particular, discuss the selection of your best model(s) using the above criteria,  $BIC$ ,  $C_P$ ,  $R_a^2$ , and  $CV_{(n)}$ . Please limit your summary to one-half page (or less).
9. Suppose we want to predict the (natural log of) the average number of caterpillar nests per tree ( $y^*$ ) for a new observed input  $\mathbf{x}^* = (1150, 20, 2, 4, 15.0, 1.1, 1.1, 6.0, 1.5, 1.3)^t$ . Use your best fitted model, chosen via  $CV(n)$ , above, to compute the predicted value, along with a (nominal) 95% prediction interval for  $y^*$ . To do this, use the `predict` function with the `interval='predict'` option, as illustrated in the last code chunk of our note chapter 2. (Incidentally, this assumes a normal linear model with an unbiased fitted model, which we might hope is approximately true; there are other methods to get a prediction interval in our case, but we have not talked about these methods yet.) Show your code/output and summarize briefly your prediction interval. In particular, comment on the validity of the nominal 95% coverage rate and the interval width.

# Bibliography

- [MR07] Jean-Michel Marin and Christian P. Robert. *Bayesian Core: A Practical Approach to Computational Bayesian Statistics*. Springer, New York, 2007.